

```
import collections
import datetime
import glob
from typing import Optional
import pathlib
import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import fluidsynth
import pretty_midi
from IPython import display
import tensorflow as tf
```

These are modules already built into python.

Data visualization libraries.

These are the relevant open-source libraries needed for this project. You'll need to download them.

✓ 0.0s

```
seed = 42
tf.random.set_seed(seed)
np.random.seed(seed) → Ensures that any random operations performed by tensorflow or numpy are consistent every time we run the code.

# Sampling rate for audio playback
sample_rate = 16000 → 16,000 Hz -> common rate for speech and some music datasets

# Downloads the maestro dataset, containing around 1200 midi files.
data_dir = pathlib.Path('data/maestro-v2_extracted/maestro-v2.0.0') → Creates a path object for the dataset's directory.

if not data_dir.exists():

    tf.keras.utils.get_file( → If the data_dir doesn't exist, downloads the MAESTRO
        'maestro-v2.0.0-midi.zip', dataset as a zip file and extracts it.
        origin='https://storage.googleapis.com/magentadata/datasets/maestro/v2.0/maestro-v2.0.0-midi.zip',
        extract=True,
        cache_dir='.', cache_subdir='data',
    )

filenames = glob.glob(str(data_dir '**/*.midi')) → Finds all the midi files within the dataset.
# print('Number of files:', len(filenames)) → Use the line below to ensure everything has
# At this stage, output should be "Number of files: 1282" run smoothly so far.
```

✓ 0.0s

```
# Practice using pretty_midi to parse a single MIDI file and inspect format of the notes
sample_file = filenames[1000]

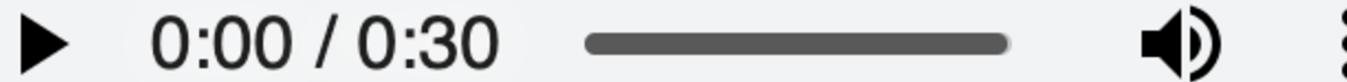
pm = pretty_midi.PrettyMIDI(sample_file)      → Takes a random file (in our case, the 1001st
                                                file) and creates an object (pm) which holds
                                                all info about MIDI data, including tracks,
                                                instruments, notes, timing, etc.

def display_audio(pm: pretty_midi.PrettyMIDI, seconds=30):
    waveform = pm.fluidsynth(fs = sample_rate)

    waveform_short = waveform[ :seconds * sample_rate] → Converts the MIDI data into an
    return display.Audio(waveform_short, rate = sample_rate)   actual digital sound file. We
                                                               specify only 30 seconds of
                                                               audio.
```

```
display_audio(pm)
```

✓ 2.0s



```
# Inspecting the instruments used in the file
print("Number of instruments:", len(pm.instruments)) → Finds the unique # of instrument
instrument = pm.instruments[0]                                tracks in a MIDI file.
instrument_name = pretty_midi.program_to_instrument_name(instrument.program)
print("Instrument name:", instrument_name)
```

✓ 0.0s

```
Number of instruments: 1
Instrument name: Acoustic Grand Piano
```

MIDI instruments are defined by an integer 0-127. This code converts the # (instrument.program) to a readable instrument name.

Here we'll take a look at the first 10 notes of our song. Enumerate ensures the index # stays with each note's info.

```
# Extract the notes from songs
for i, note in enumerate(instrument.notes[:10]):
    note_name = pretty_midi.note_number_to_name(note.pitch) Like program_to_instrument_name,
    duration = note.end - note.start                         this function converts the MIDI
    print(f"{i}: pitch = {note.pitch}, note_name = {note_name}, integer for pitch to a human
          |   |   f' duration = {duration:.4f}'                  readable name.
# For the first 10 notes in the song, output is the pitch, note name, and duration of each note
# Pitch is perceptual quality of sound as a MIDI note number
```

✓ 0.0s

```
0: pitch = 49, note_name = C#3, duration = 0.6198
1: pitch = 77, note_name = F5, duration = 1.0130
2: pitch = 73, note_name = C#5, duration = 0.3438
3: pitch = 56, note_name = G#3, duration = 0.6341
4: pitch = 56, note_name = G#3, duration = 0.0755
5: pitch = 56, note_name = G#3, duration = 0.0872
6: pitch = 65, note_name = F4, duration = 1.2135
7: pitch = 61, note_name = C#4, duration = 1.2057
8: pitch = 68, note_name = G#4, duration = 1.5430
9: pitch = 56, note_name = G#3, duration = 0.4180
```

```
#Extract notes from a MIDI file
def midi_to_notes(midi_file: str):
    pm = pretty_midi.PrettyMIDI(midi_file)
    instrument = pm.instruments[0]
    notes = collections.defaultdict(list) →
    #Sort the notes by start time
    sorted_notes = sorted(instrument.notes, key=lambda note: note.start)
    prev_start = sorted_notes[0].start →
    for note in sorted_notes:
        start = note.start
        end = note.end
        notes["pitch"].append(note.pitch)
        notes["start"].append(note.start) →
        notes["end"].append(end)
        notes["step"].append(start - prev_start)
        notes["duration"].append(end - start)
        prev_start = start
    return pd.DataFrame({name: np.array(value) for name, value in notes.items()})
```

Notes gives us a default dictionary where each key has an empty list, allowing us to append values.

Sorts our notes based on start time. Prev\_start is initialized to the start time of the first note. We'll later use this to calculate "step".

Now, each note in our song has a unique value for pitch, duration, and step.

```
raw_notes = midi_to_notes(sample_file)
```

```
raw_notes
```

✓ 0.1s

This is a dataframe of  
the pitch, step, and  
duration of every note  
in the MIDI file!

	<b>pitch</b>	<b>start</b>	<b>end</b>	<b>step</b>	<b>duration</b>
0	77	0.976562	1.989583	0.000000	1.013021
1	49	1.006510	1.626302	0.029948	0.619792
2	56	1.636719	2.270833	0.630208	0.634115
3	73	1.916667	2.260417	0.279948	0.343750
4	68	2.140625	3.683594	0.223958	1.542969
...	...	...	...	...	...
1496	61	273.015625	279.102865	0.015625	6.087240
1497	37	273.016927	278.799479	0.001302	5.782552
1498	56	273.022135	279.074219	0.005208	6.052083
1499	38	282.093750	282.264323	9.071615	0.170573
1500	45	282.264323	282.278646	0.170573	0.014323

```
get_note_names = np.vectorize(pretty_midi.note_number_to_name)  
sample_note_names = get_note_names(raw_notes['pitch'])  
raw_notes = raw_notes.assign(pitch = sample_note_names)
```

This function converts pitch # to  
a readable name.

1501 rows x 5 columns

```

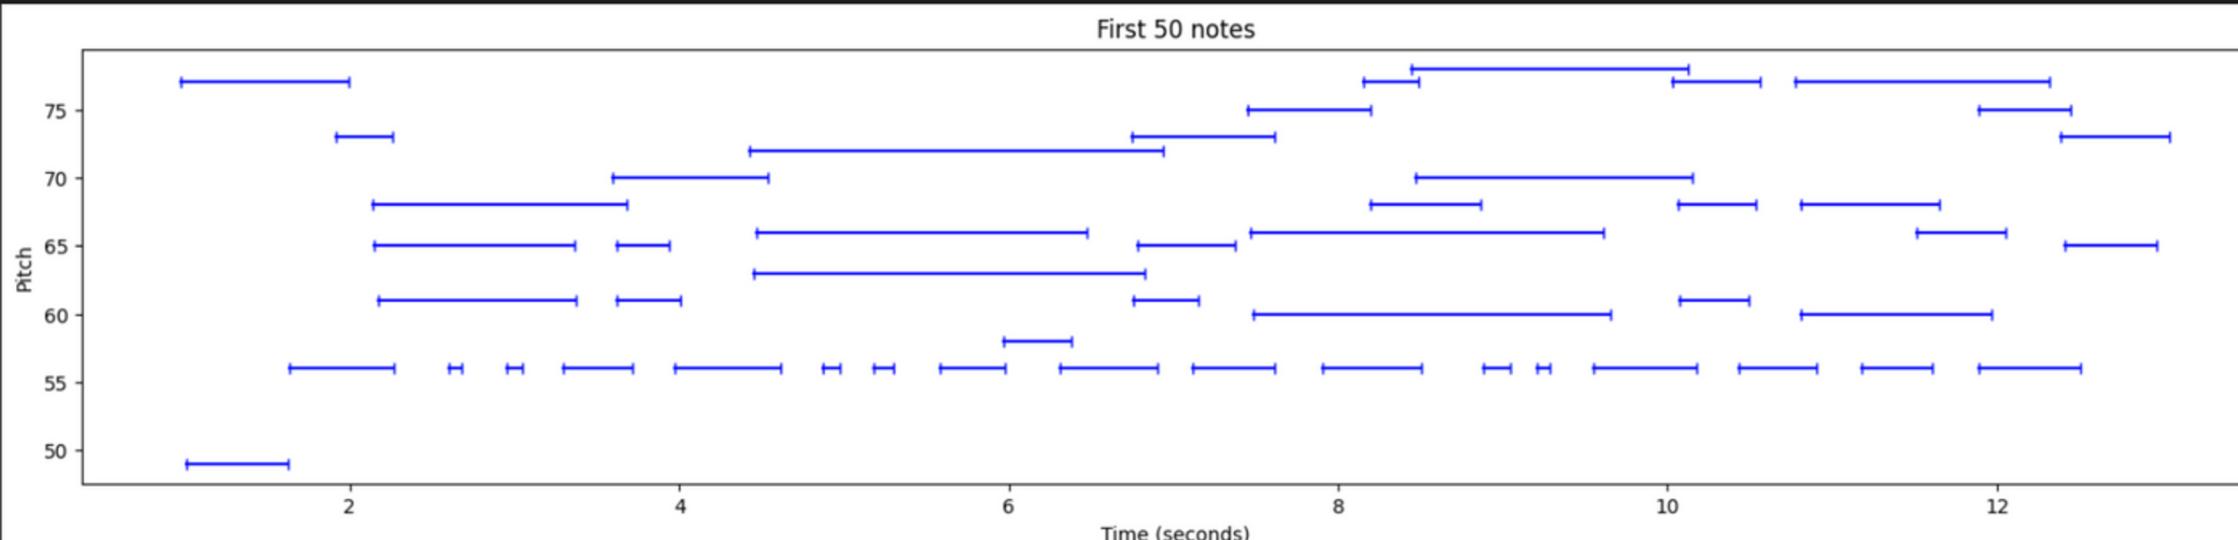
def plot_piano_roll(notes: pd.DataFrame, count: Optional[int] = None):
    if count:
        title = f'First {count} notes' → If the optional count is provided, then
    else:                                only that # of notes will be displayed.
        title = f'Whole track'
        count = len(notes['pitch'])
    plt.figure(figsize=(20, 4))
    plot_pitch = np.stack([notes['pitch'], notes['pitch']], axis=0)
    plot_start_stop = np.stack([notes['start'], notes['end']], axis=0) → Stacks the columns twice
    plt.plot(                                along axis 0, creating a 2d
        plot_start_stop[:, :count], plot_pitch[:, :count], color="blue", marker="|")
    plt.xlabel('Time (seconds)')
    plt.ylabel('Pitch')
    _ = plt.title(title)

```

plot\_piano\_roll(raw\_notes, count=50)

✓ 0.2s

Python



Here's an example of the first 50 notes, plotted with respect to duration and pitch!

Stacks the columns twice along axis 0, creating a 2d array. The first copy is for the start time, the second is for the end time.

```
#check distribution of each note variable
def plot_distributions(notes: pd.DataFrame, drop_percentile=2.5):
    plt.figure(figsize=[15, 5])
    plt.subplot(1, 3, 1)
    sns.histplot(notes, x="pitch", bins=20)

    plt.subplot(1, 3, 2)
    max_step = np.percentile(notes['step'], 100 - drop_percentile)
    sns.histplot(notes, x="step", bins=np.linspace(0, max_step, 21))

    plt.subplot(1, 3, 3)
    max_duration = np.percentile(notes['duration'], 100 - drop_percentile)
    sns.histplot(notes, x="duration", bins=np.linspace(0, max_duration, 21))
```

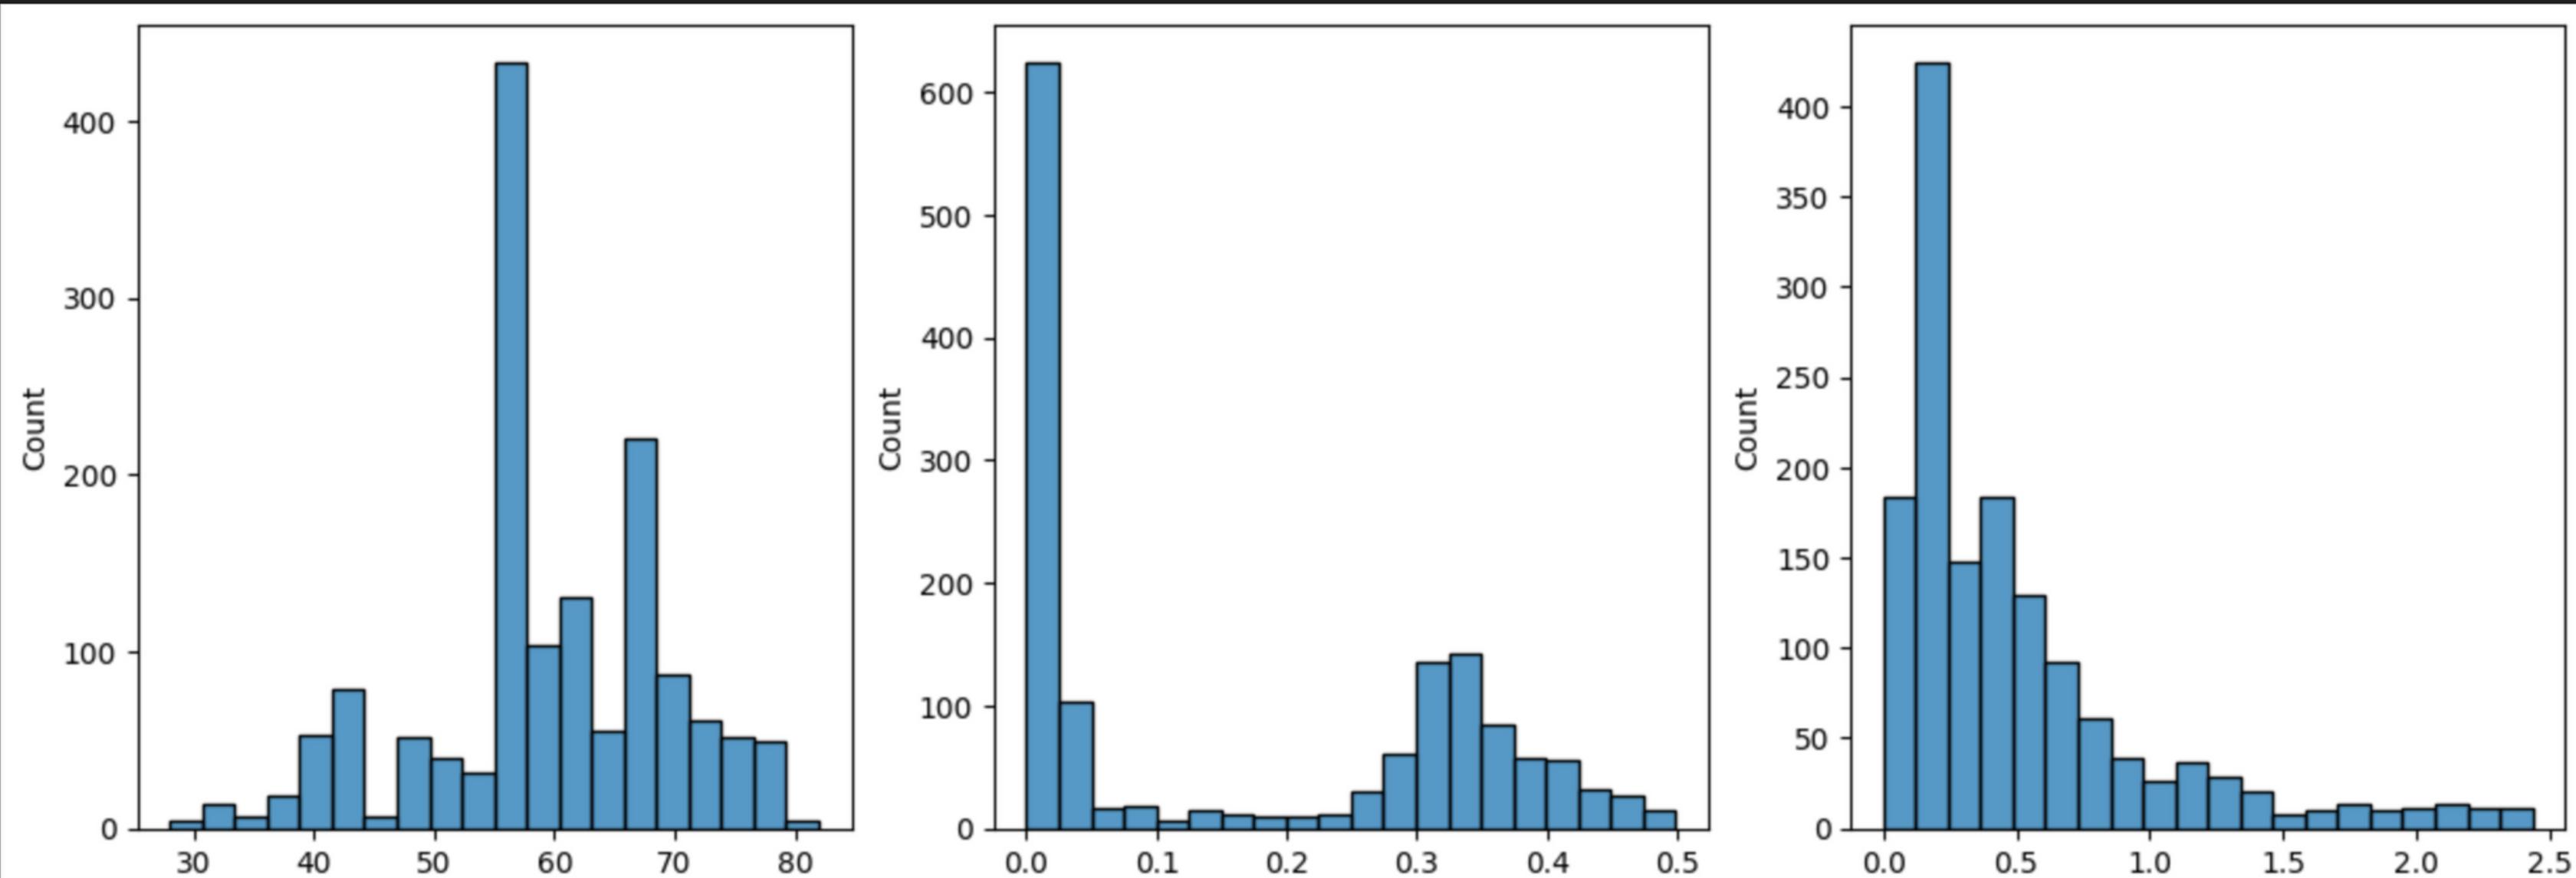
Drop\_percentile is used to drop outliers

Creating three histograms, at the 97.5th percentile, for pitch, step, and duration of our chosen song sample

✓ 0.0s

plot\_distributions(raw\_notes)

✓ 0.7s



```
#create your own MIDI file
def notes_to_midi(
    notes: pd.DataFrame, out_file: str, instrument_name: str, velocity: int=100
):
    pm = pretty_midi.PrettyMIDI()
    instrument = pretty_midi.Instrument(
        program = pretty_midi.instrument_name_to_program(instrument_name))
    Creates an instrument object that
    will be added to the MIDI file

    prev_start = 0
    for i, note in notes.iterrows():
        start = float(prev_start + note["step"])
        end = float(start + note["duration"])
        .iterrows iterates through each
        row of the notes dataframe.
        note = pretty_midi.Note(
            velocity = velocity, pitch = int(note["pitch"]), start = start, end = end
        )
        instrument.notes.append(note)
        prev_start = start
        The created note ia appended to the
        notes list of the instrument object

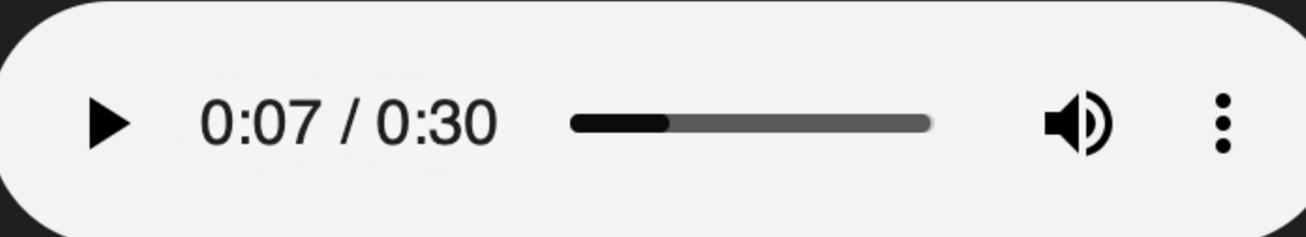
    pm.instruments.append(instrument)
    Now that all notes are added, instrument
    is appended to the pm.instruments list
    pm.write(out_file)
    return pm
```

```
example = "example.midi"
example_pm = notes_to_midi(raw_notes, out_file = example, instrument_name = instrument_name)
```

✓ 0.1s

```
#play the generated MIDI file
display_audio(example_pm)
```

✓ 1.6s



0:07 / 0:30

```
#creating the training dataset
#extract notes from the MIDI files, starting with a small number of files
num_files = 5      → extracts data from a set of MIDI files and creates a dataset of
all_notes = []
for f in filenames[:num_files]:
```

```
    notes = midi_to_notes(f)      → Slices the filenames list to only include the first 5 MIDI file
    all_notes.append(notes)
```

```
all_notes = pd.concat(all_notes)
```

paths based on num\_files. The notes for each file are appended to all\_notes